

ServiceMeshCon

# How (and Why) Google Manages Millions of Sidecars Without Breaking Everything

*Sven Mawson (sven@google.com)*

# The Sidecar



# Why a Sidecar?



- Prohibitive per-language expense
  - HTTP and RPC client and server libraries
  - Service Management (Networking, Security, Telemetry)
  - API Management
  - Other complex client libraries (e.g. Data Storage, Distributed File Systems, Messaging, Locking, ...)

# Why a Sidecar?

- Ongoing per-language divergence
  - Default settings
  - Bugs (bug-for-bug compatibility)
  - Performance profiles

# Why a Sidecar?

- Difficulties in rolling out fleet-wide changes
  - How many processes have bug X?
  - Rebuild, re-qualify, redeploy
  - What if it is a critical fix?

# But why a *Sidecar*!



- Rejected Alternatives
  - Foreign Function Interfaces (e.g. SWIG)
  - DLLs
  - Remote Services

Roll out a sidecar to millions of processes, keep it up to date, and don't break anything.





# Challenge Accepted!

# Enabling Sidecars is Scary!



- Start with opt-in for those with something to gain
  - New languages (e.g. golang)
  - New functionality (e.g. Api Management)

# Enabling Sidecars is Scary!



Make It  
Fast &  
Cheap



# Enabling Sidecars is Scary!



- Make it fast, make it cheap
  - Multiply latency savings by # of hops
  - Multiply resource savings by # of sidecars
  - Never stop investing in performance improvements

# Enabling Sidecars is Scary!



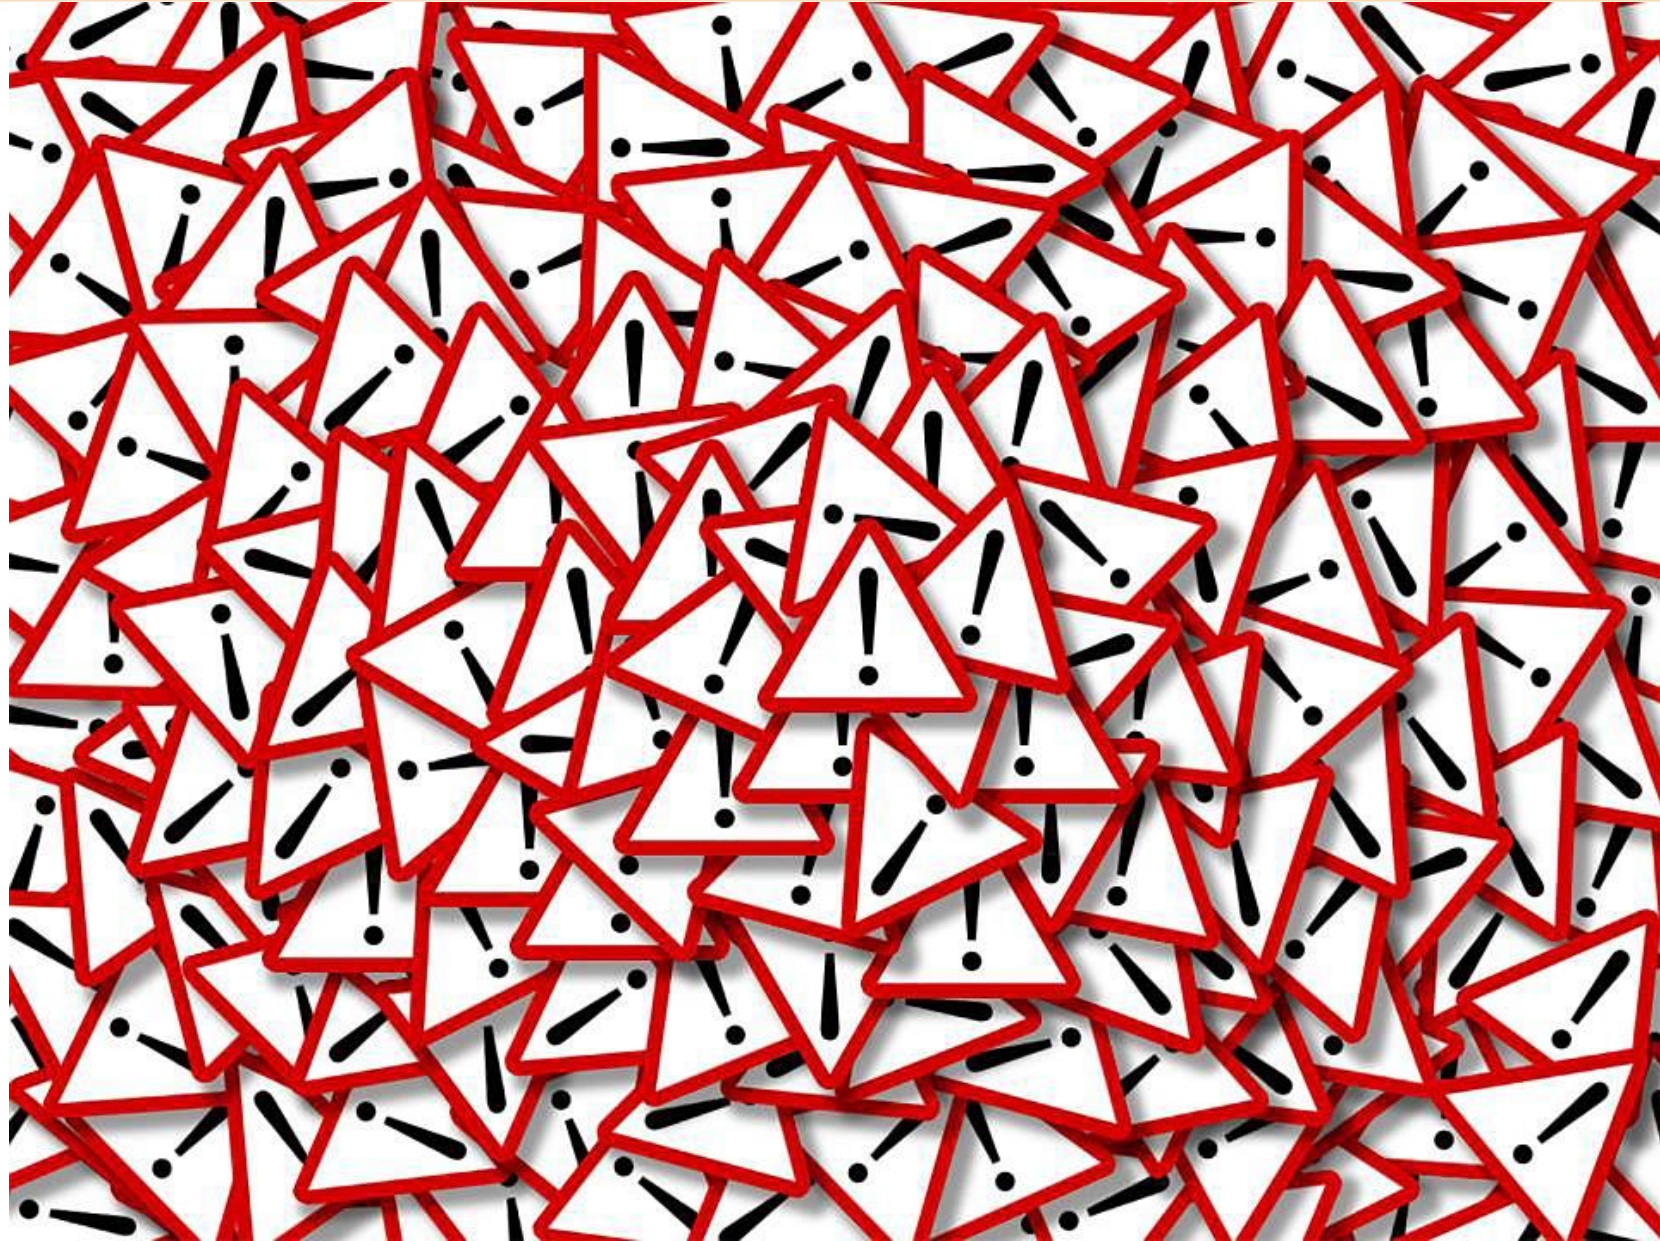
- Proactively identify those most affected
  - Anyone close to a resource limit already
  - Anyone with a business-critical job



# Enabling Sidecars is Scary!



Warn  
Your  
Users!



# Enabling Sidecars is Scary!



- Warn them!
  - And warn them again. And again.
  - And again.



# Enabling Sidecars is Scary!



Take It  
Sloooooow!





# Enabling Sidecars is Scary!



- Enable it by default sloooooooooowly
  - Cluster by cluster, job by job
  - Make opt-out easy, follow up later

# What About Upgrades?

# Upgrading is also Scary!



- Two tracks: dev and stable
  - dev is weekly, stable is monthly

# Upgrading is also Scary!



Keep  
It  
Simple!



# Upgrading is also Scary!



- Two modes: automatic and bound
  - Automatic happens on restarts (~monthly)
  - Bound mode sets version = latest when built
  - Bound mode requires agreement to:
    - Update regularly (at least every 2 weeks)
    - Update quickly for critical fixes



# Upgrading is also Scary!



# Still Sloooooow!

# Upgrading is also Scary!



- Roll out updates slowly!
  - Qualify new releases with critical users
  - Start with canaries (automatic + bound)
  - Slowly ramp to 100%
  - Take your time, especially early on!



# Upgrading is also Scary!



Don't  
Make  
Them  
Grumpy





# Upgrading is also Scary!



- Make rollback self-serve
  - Allow users to roll back to a specific supported version
  - Keep a dashboard of how many users rolled back
  - Be proactive and fix their issues ASAP
  - Roll the whole thing back if too many users broke

OK, So What Did We  
Learn?

# Lessons Learned



- Invest in debugging tools
  - Whose fault is it?
  - How often is it happening?

# Lessons Learned



- Application <-> Sidecar comms need to be rock solid
  - Tried an experimental shared memory channel
  - Ended up with Unix Domain Socket: slower but much more reliable

# Lessons Learned



- You're going to take a performance hit
  - Batch chatty connections
  - Keep investing in performance!

# Lessons Learned



- Avoid changes in default behavior
  - Dynamically load functionality on demand
  - Guard new functionality with opt-in and opt-out
  - Be especially careful of newly opened connections!

- Don't get fancy
  - Share fate between application and sidecar
  - No hot reloading, pick up new versions on restart
  - Avoid one-offs
  - Make it easy to run the sidecar during development
  - Make it easy to run the sidecar in tests

It is Totally Worth  
It!



# Questions?